Global Big Data Conference

GLOBAL
ARTIFICIAL INTELLIGENCE
CONFERENCE

BOSTON

OCT 1st, 2nd, 3rd 2019

Boston Convention & Exhibition Center

www.globalbigdataconference.com

Twitter : @bigdataconf
#GAIC

Financial Loan Risk Analysis Using Keras Deep Learning

Marty Lurie, marty@cloudera.com

The business question: Can we use AI/Deep Learning to predict if a loan will be repaid? Specifically, can we predict if there will be foreclosure costs associated with a loan based on the features of the loan?

Breaking news September 24, 2019… (today is October 3, 2019)

THE WALL STREET JOURNAL.

Subscribe | Sign In

U.S. Edition ▾ | September 24, 2019 | Print Edition | Video

Home World U.S. Politics Economy Business Tech Markets Opinion Life & Arts Real Estate WSJ. Magazine Search

BREAKING NEWS House Speaker Nancy Pelosi to announce formal impeachment inquiry of President Trump, says a person familiar with the matter

SHARE

TEXT

MARKETS

Freddie Mac Tests Underwriting Software That Could Boost Mortgage Approvals

Regulator recently met with housing-finance giant and the fintech firm behind the software

Agenda

- Statistics, Machine Learning, and AI, Oh My!
- Why Fannie Mae Loan Analysis?
- Architecture for ML/Deep Learning
- Data Engineering
- Logistic Regression
- Keras Neural Network
- Results and what's next
- Here are the whole works: https://github.com/git4impatient/fanniemae

Statistics, Machine Learning, and AI, Oh My!

- ***Statistics*** is the discipline that concerns the collection, organization, displaying, analysis, interpretation and presentation of data.[1][2][3]
- ***Machine learning (ML)*** is the scientific study of algorithms and statistical models that computer systems use to perform a specific task without using explicit instructions, relying on patterns and inference instead. It is seen as a subset of artificial intelligence.
- In computer science, ***artificial intelligence (AI)***, sometimes called machine intelligence, is intelligence demonstrated by machines, in contrast to the natural intelligence displayed by humans.
- All definitions from https://en.wikipedia.org

Anyone ever try to change a definition in wikipedia?

Why Fannie Mae Loan Analysis?

- The Federal National Mortgage Association (FNMA), commonly known as Fannie Mae, is a United States government-sponsored enterprise (GSE) and, since 1968, a publicly traded company. Founded in 1938 during the Great Depression as part of the New Deal,[2] the corporation's purpose is to expand the secondary mortgage market by securitizing mortgage loans in the form of mortgage-backed securities (MBS),[3] *Source: wikipedia.org*
- This talk is based on data from: The Home Affordable Refinance Program (HARP) is a federal refinance program targeting underwater homeowners. *https://www.hsh.com/finance/refinance/what-is-harp-do-i-qualify-for-a-harp-loan.html*

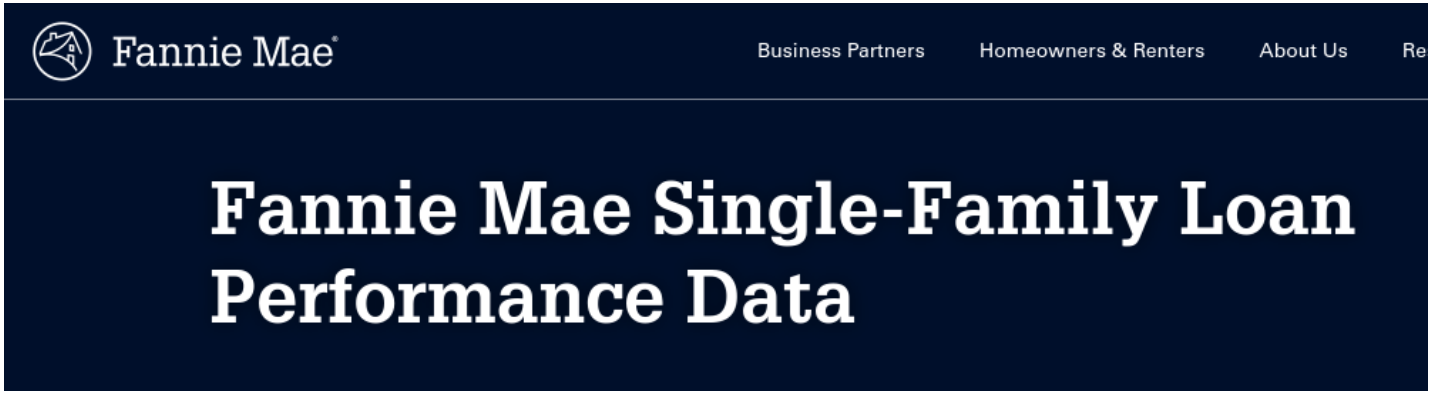

https://www.fanniemae.com/portal/funding-the-market/data/loan-performance-data.html

Architecture for ML/Deep Learning

Projects/Software Used:

- Apache Spark
- Apache Impala
- HUE
- Apache Hadoop
- Tensorflow
- Keras
- Pandas
- Numpy
- Seaborn
- Matplotlib
- CDSW

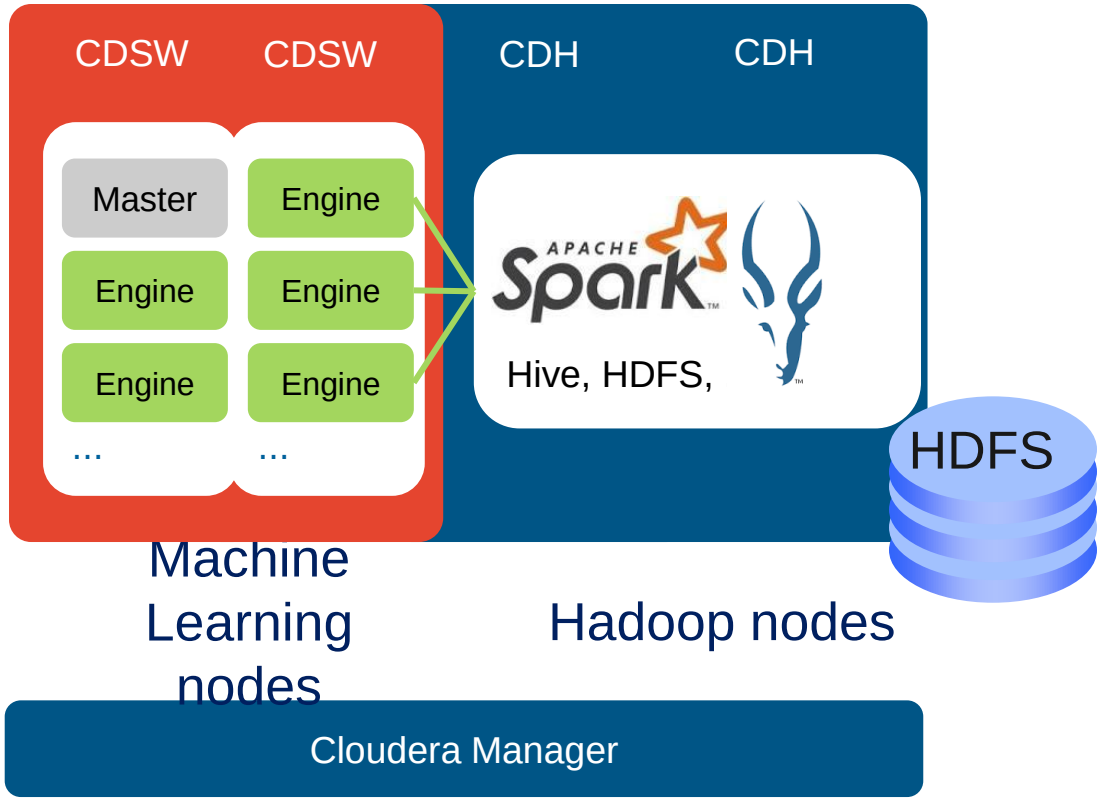

© Cloudera, Inc. All rights reserved.

Data Engineering

- "80% of analysts' time is spent discovering and preparing data" *Harvard Business Review*
- Source data: Fannie Mae HARP
  - Acquisition file: loan origination details

```
[marty@gromit HARP]$ head Acquisition_HARP.txt.sample
100001565398|R|QUICKEN LOANS INC.|4.375|153000|360|07/2012|09/2012|138|138|1||700|N|R|SF|1|I|FL|347|
100003305358|R|OTHER|4.75|342000|360|04/2009|06/2009|94|94|2||734|N|R|SF|1|P|NY|117|18|FRM|733|1|N
100004116882|R|PNC BANK, N.A.|3.875|93000|180|06/2014|08/2014|105|105|1||526|N|R|SF|1|P|IL|608||FRM|
100006858918|R|CITIMORTGAGE, INC.|4.125|105000|360|12/2012|03/2013|81|81|2||570|N|R|SF|1|P|CA|953||F
```

  - Performance file: payment status, repeating rows

```
[marty@gromit HARP]$ head Performance_HARP.txt.sample
100001565398|08/01/2012|QUICKEN LOANS INC.|4.375||0|360|359|08/2042|29460|0|N|||||||||||||||||N
100001565398|09/01/2012|OTHER|4.375||1|359|359|08/2042|29460|0|N|||||||||||||||||Y
100001565398|10/01/2012||4.375||2|358|358|08/2042|29460|0|N|||||||||||||||||N
100001565398|11/01/2012||4.375||3|357|357|08/2042|29460|0|N|||||||||||||||||N
100001565398|12/01/2012||4.375||4|356|356|08/2042|29460|0|N|||||||||||||||||N
```

Data Engineering – Create Tables in Impala

```
use fanniemae;
drop table if exists loan_perf;
create external table loan_perf(
loan_identifier string,
monthly_reporting_period string,
servicer_name string,
current_interest_rate decimal ( 14,10 ) yada yada yada
… lines deleted …
foreclosure_principal_write_off_amount decimal ( 11,2 ),
servicing_activity_indicator string
)
row format delimited fields terminated by '|'
stored as textfile location '/user/marty/fanniemae/perf'
;

select count(*) from loan_perf;
select * from loan_perf limit 1;
```

Data Engineering – SQL w/Impala Example

- SparkML, Tensorflow require numeric variables, best if scaled to values 0 to 1

```
drop table if exists loanacq_sqlnormed_p ;
create table loanacq_sqlnormed_p stored as parquet as
select loan_identifier,
case origination_channel
when 'R' then .1
when 'C' then .2
when 'B' then .3
end channel,
seller_name,
original_interest_rate/7.75 intrate,
original_upb/1402000.0 loanamt,
original_loan_to_value/ 97 loan2val,
number_of_borrowers/6 numborrowers,
--original_debt_to_income_ratio/64 debt2income,
borrower_credit_score_at_origination/842 creditscore,
property_state,
origination_date
from loan_acquisition;
```

Data Engineering – pySpark Example

- SparkML, Tensorflow require numeric variables, best if scaled to values 0 to 1

```
# need to convert from text field to numeric
# common requirement when using sparkML
from pyspark.ml.feature import StringIndexer
# this will convert each unique
# string into a numeric
indexer = \
StringIndexer(inputCol="property_state", \
outputCol="loc_state")
indexed = indexer.fit(lndf).transform(lndf)
indexed.show(5)
```

Data Engineering – Know Your Data
seaborn pairplot

```
g = sns.pairplot(pdsdf,  hue="seller_name" )
```

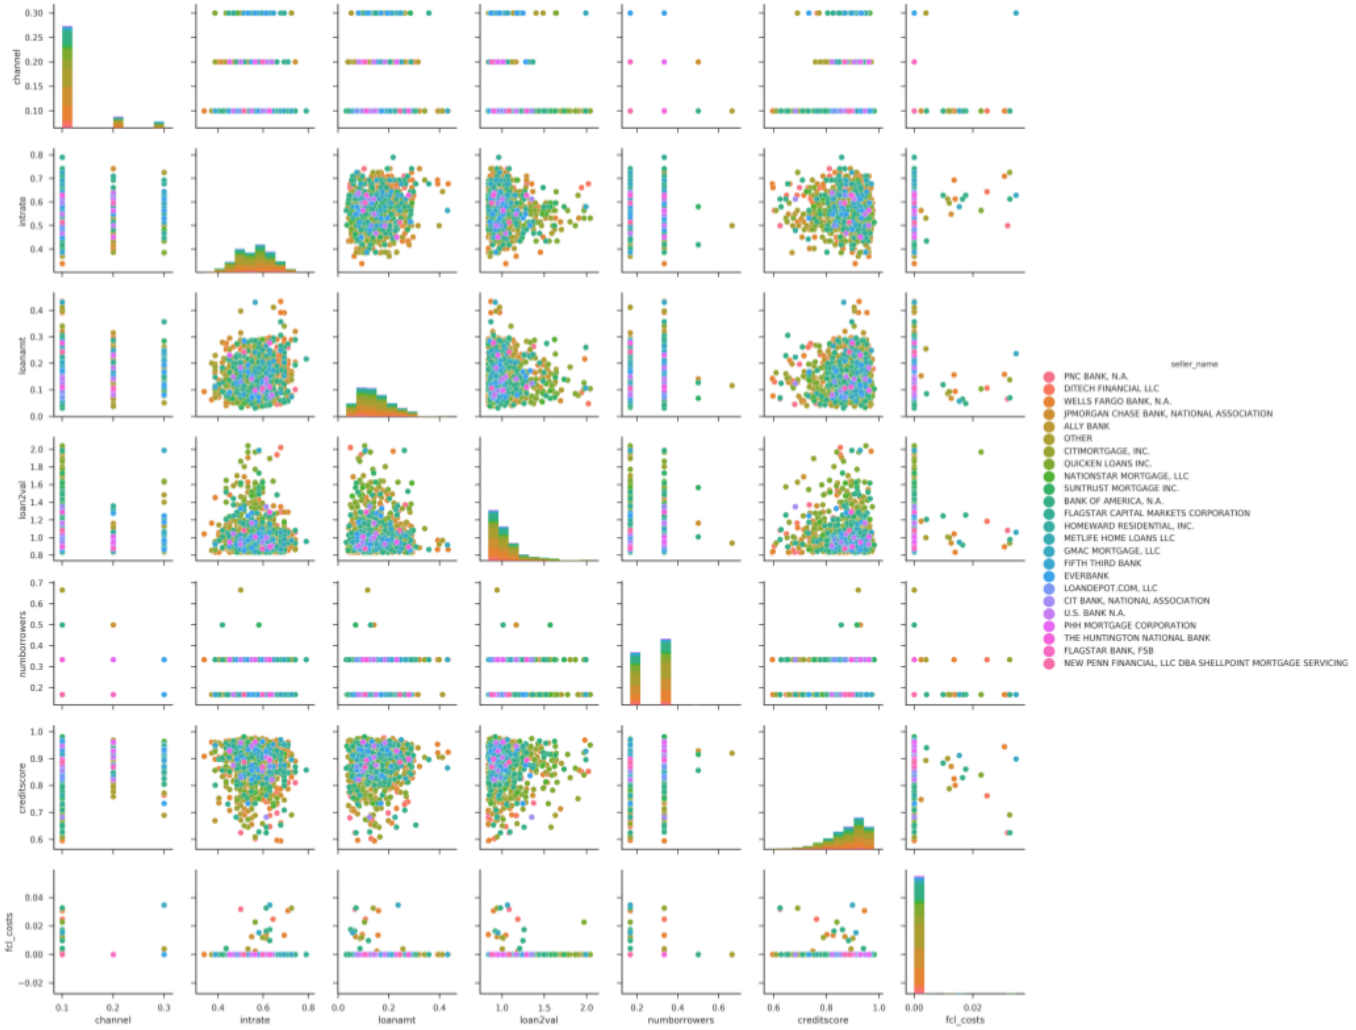

Data Engineering – Know Your Data seaborn heatmap

```
sns.heatmap(kerasinputpsdf.corr(), annot=True)
```

```
<matplotlib.axes._subplots.AxesSubplot at 0x7f6622cda310>
```

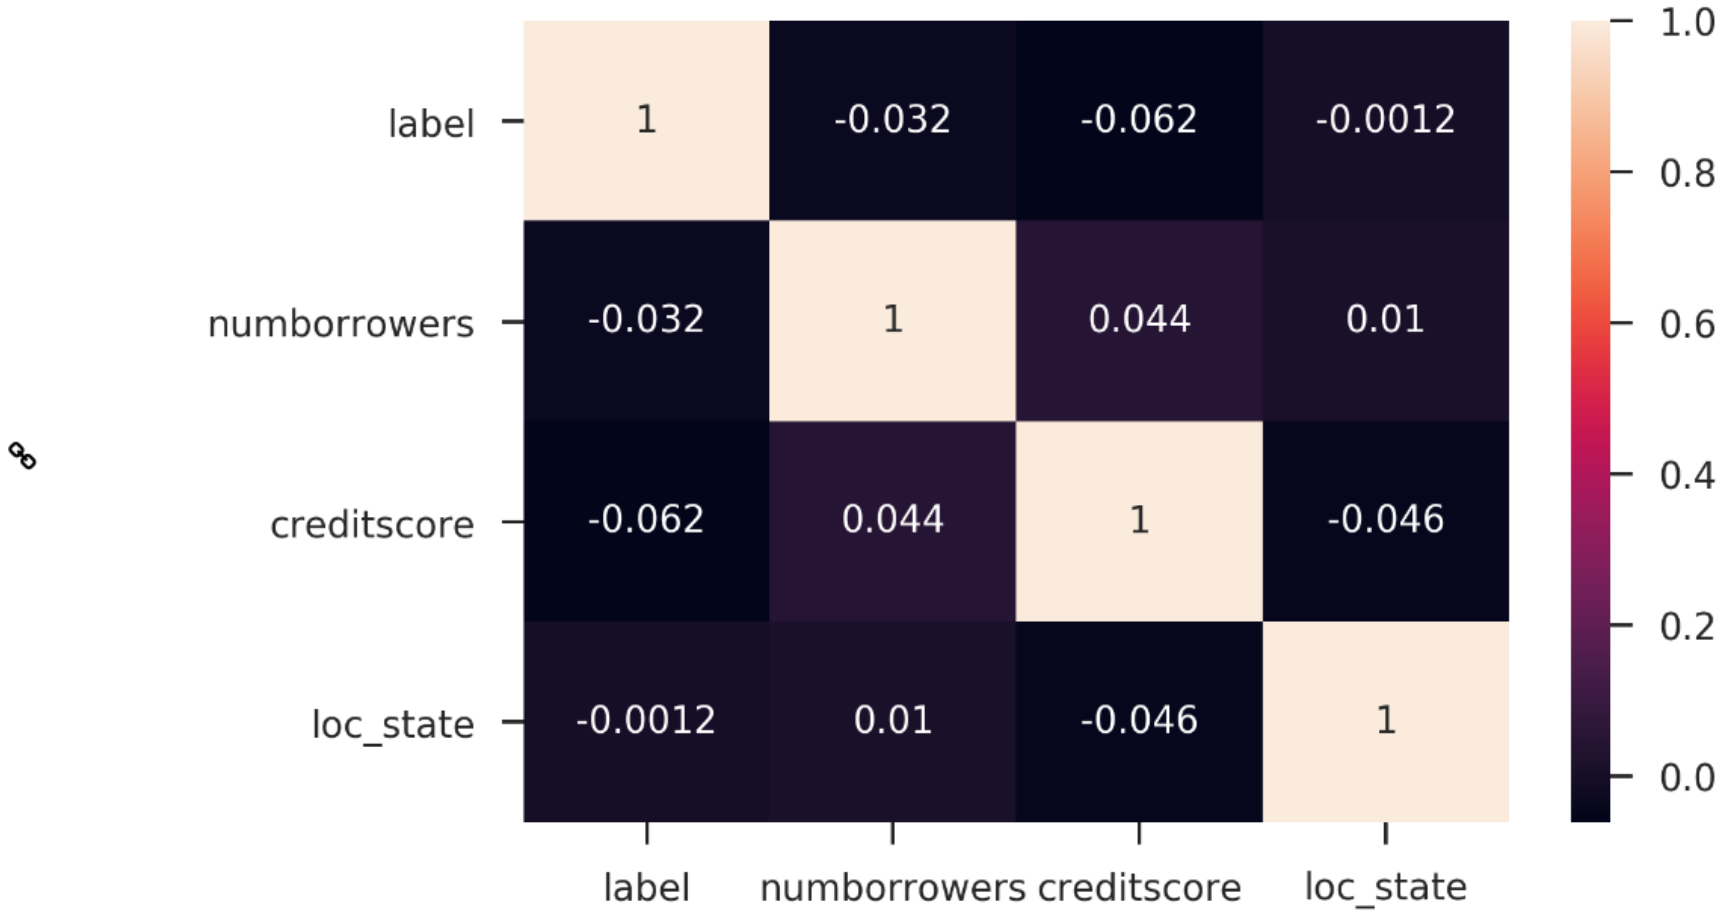

Data Engineering – Know Your Data
Example with stronger correlations

http://www.codeheroku.com/post.html?name=Introduction%20to%20Exploratory%20Data%20Analysis%20(EDA)

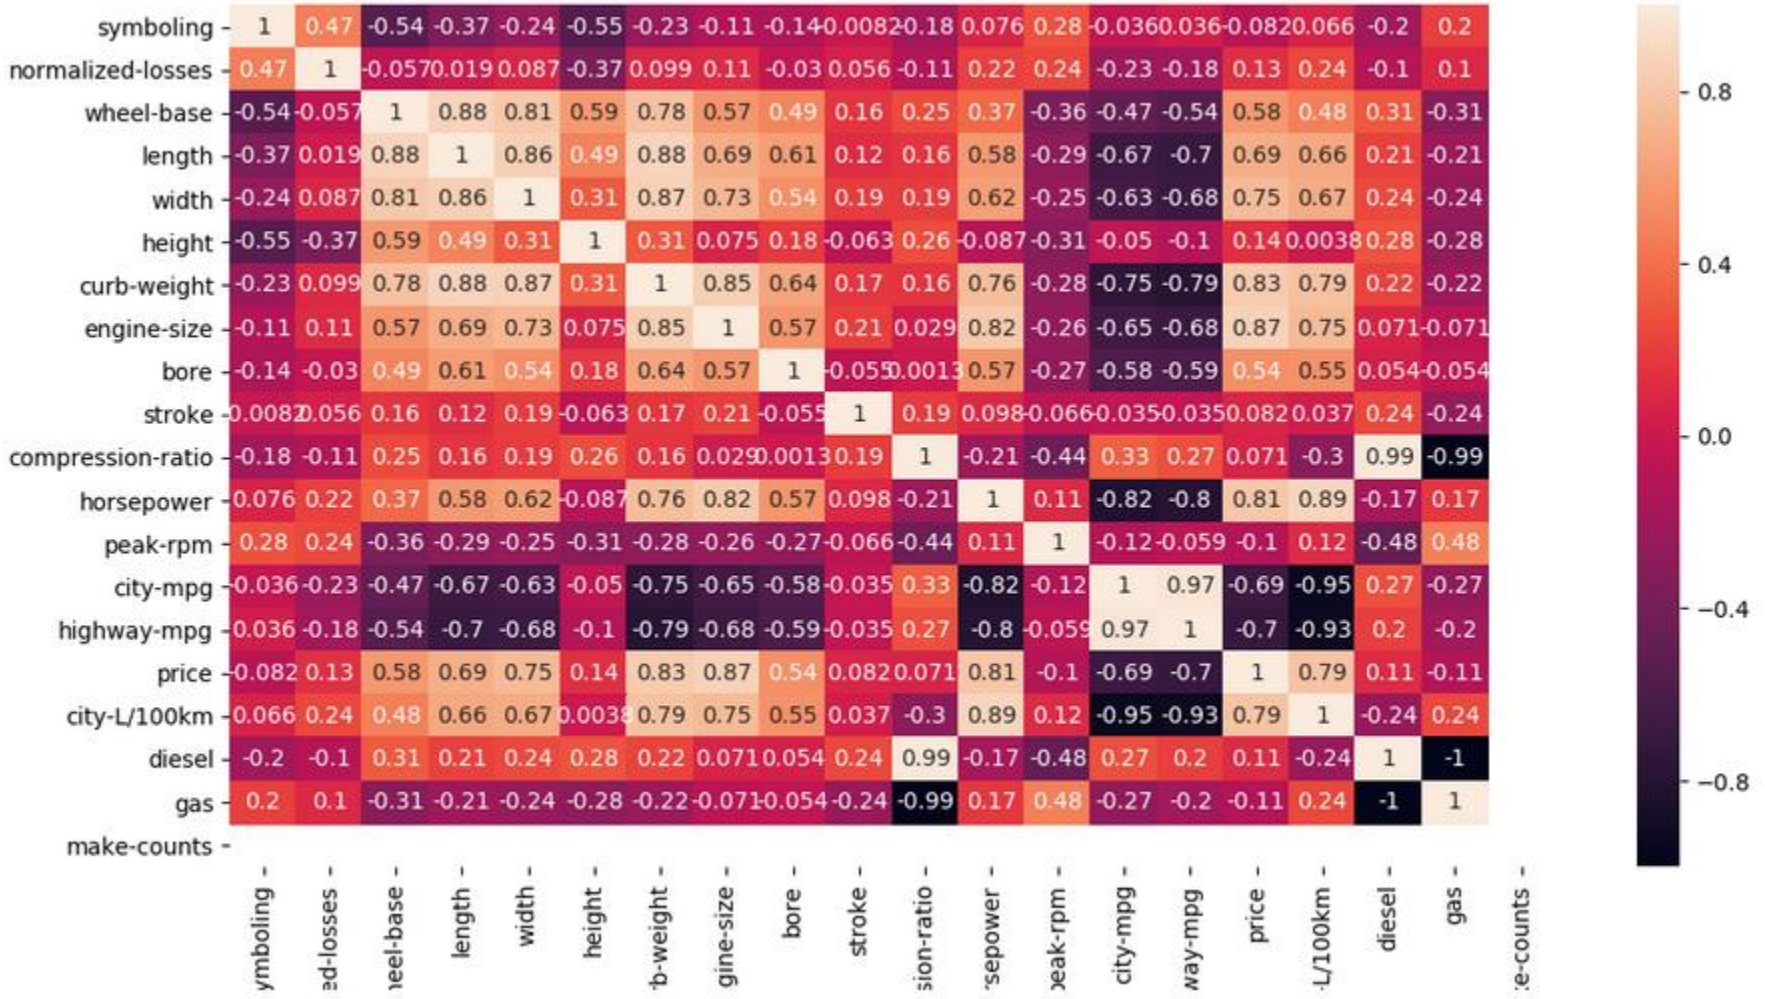

Logistic Regression (Classification)

- *"Logistic regression is the appropriate regression analysis to conduct when the dependent variable is dichotomous (binary)."*
  https://www.statisticssolutions.com/what-is-logistic-regression/
- Two different Fannie Mae loan distributions, one with Foreclosure Costs one without Foreclosure costs:

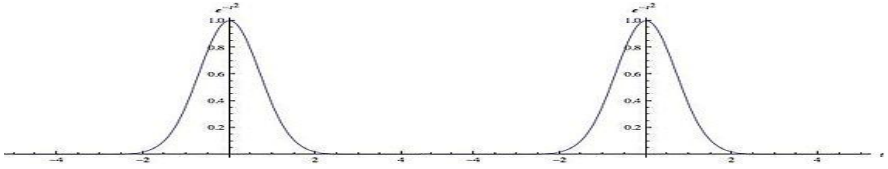

- Does this help?

$$\text{logit}(p) = \log\left(\frac{p(y=1)}{1-(p=1)}\right) = \beta_0 + \beta_1 \cdot x_{i2} + \beta_2 \cdot x_{i2} + \_ + \beta_p \cdot x_{in}$$

for i = 1…n .

Logistic Regression

- pySpark Logistic Regression Inputs:
  - spark dataframe with: label and features

LABEL	FEATURES
0 OR 1 for foreclosure costs	[income, loan value, interest rate, etc]

```
# Create a LogisticRegression instance
lr = LogisticRegression(maxIter=10, regParam=0.01)
# Print out the parameters, documentation, and any default values.
print("LogisticRegression parameters:\n" + lr.explainParams() + "\n")
# Learn a LogisticRegression model. This uses the parameters stored in lr.
#  “output” is a lousy name for the input-dataframe, sorry ☹
model1 = lr.fit(output)
```

Keras Neural Network

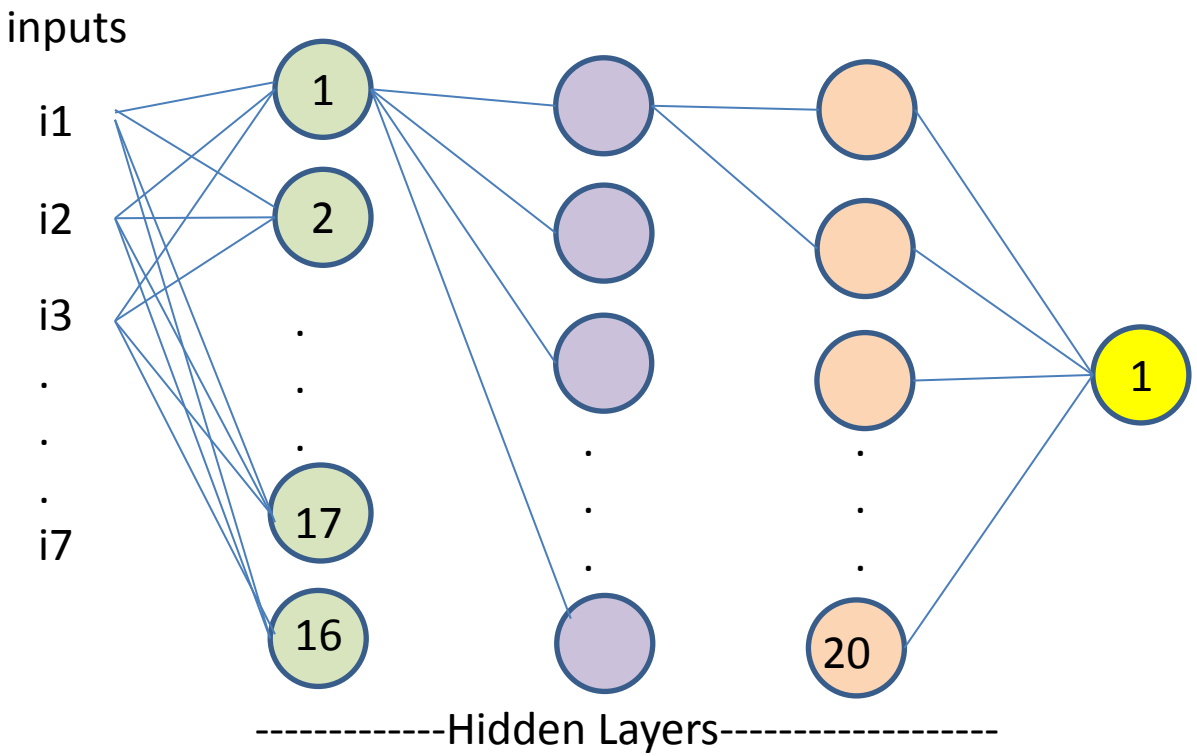

Single output neuron

Not all connections show

Keras Neural Network

```
classifier = Sequential()
#First Hidden Layer
classifier.add(Dense(16, activation='relu', kernel_initializer='random_normal', input_dim=7))
#Second  Hidden Layer
classifier.add(Dense(20, activation='relu', kernel_initializer='random_normal'))
#Output Layer
classifier.add(Dense(1, activation='sigmoid', kernel_initializer='random_normal'))
#Compiling the neural network
classifier.compile(optimizer ='adam',loss='binary_crossentropy', metrics =['accuracy'])
```

```
classifier.summary()
```

```
_________________________________________________________________
Layer (type)                 Output Shape              Param #
=================================================================
dense_1 (Dense)              (None, 16)                128
_________________________________________________________________
dense_2 (Dense)              (None, 20)                340
_________________________________________________________________
dense_3 (Dense)              (None, 1)                 21
=================================================================
Total params: 489
Trainable params: 489
Non-trainable params: 0
```

Results

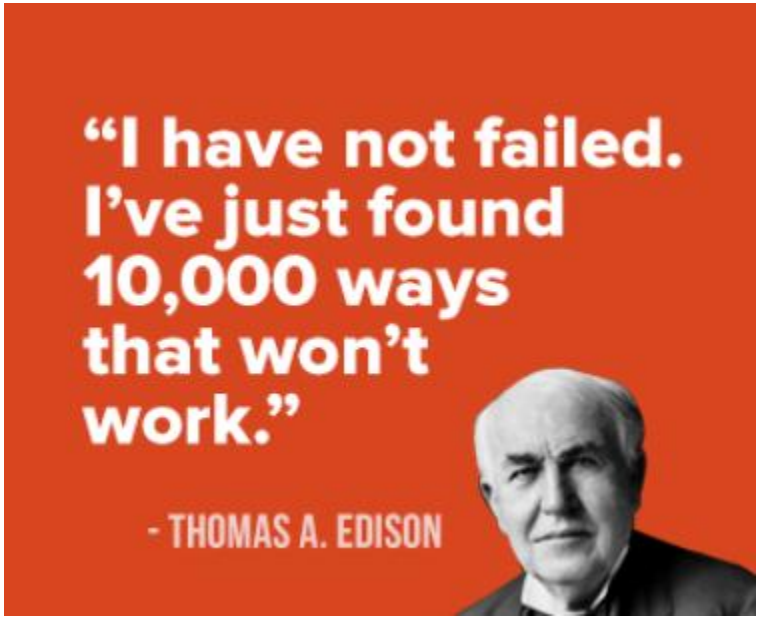
“I have not failed.
I’ve just found
10,000 ways
that won’t
work.”
- THOMAS A. EDISON

Results – Confusion Matrix

- Baseline Probability, incidence of foreclosure costs: 0.0213
- Logistic Regression, areaUnderROC: 0.738:

All Data (easy case)	Predict 0	Predict 1
Expect 0	502924	0
Expect 1	10942	0

Non-Diagonal is error

Diagonal is good

- Keras Deep Learni… Prob of…osts 2183/100571=0.0217

Training Data	Predi…	Predict 1
Expect 0	402353	0
Expect 1	69	8671

Testing Data	Predic…	Predict 1
Expect 0	1005…	
Expect 1	19…	3

Looked really good BUT the value of foreclosure costs was provided in the input matrix!

Results – Confusion Matrix

- Baseline Probability, incidence of foreclosure costs: 0.0213
- Logistic Regression, areaUnderROC: 0.738:

All Data (easy case)	Predict 0	Predict 1
Expect 0	502924	0
Expect 1	10942	0

Non-Diagonal is error

Diagonal is good

- Keras Deep Learning

Testing Data	Predict 0	Predict 1
Expect 0	97586	2985
Expect 1	1917	285

Wow we have a model!
Are we done?
Not Yet…

What's Next?

- We have a model, how do we use it? RestAPI

```
[marty@gromit fanniemae]$ curl -H "Content-Type: application/json" -X POST http://cdsw2.lurie.biz/api/altus-ds-1/models/call-model -d
'{"accessKey":"mz7j447kq5q22o1vf4877fpdm6gl4n70","request":{"channel":0.1,"intrate":0.4,"loanamt":0.18,"loan2val":0.9,"numborrowers":0.33,"creditscore":0.91}}'
{
  "success": true,
  "response": "False"
}
```

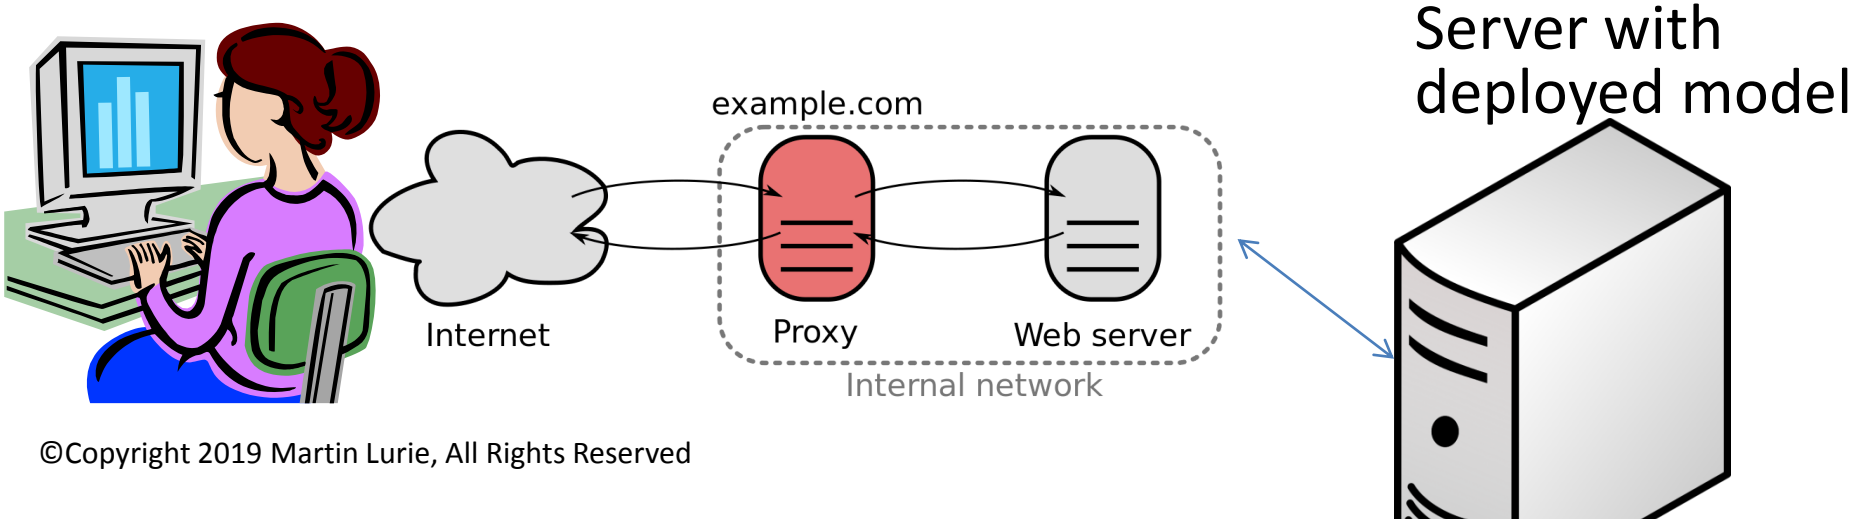

What's Next?

- Add more features, enrich data with external datasets
- Experiment with Hyperparameters: layers, neurons etc
- TensorBoard for model validation
- Model deployment – restAPI
- Full FannieMae dataset
  - HARP is subset of data that fits on my laptop
  - Some columns ignored that may improve model

Thank You!
Questions?

marty@cloudera.com